



Yang, Bo

Machine Learning Engineer

- +1-415-420-9693
- https://yang-bo.com
- yang-bo@yang-bo.com
- 955091
- Atry
- yang-bo

About me


I am a Machine Learning Engineer focused on architecting and operationalizing reliable, high-leverage ML and language model platforms end-to-end.

I blend deep systems engineering (performance, reliability, observability) with functional programming to create concise, extensible, and correct abstractions for data, model, and multi-agent workflows.

A long-time open-source maintainer (3000+ GitHub stars across frameworks for neural networks, reactive UI, micro-services, and DSLs), I emphasize automation, clarity of design, and sustainable velocity over ad-hoc complexity.

I bring a product mindset, rigorous engineering standards, and a bias toward serverless scalability, safety, and measurability in every ML system I help shape.

Skills



Experience

- 2024- Tacnode Machine Learning Engineer
Built a serverless platform for LLM fine-tuning, test-time training, and inference, integrated with Discord and Hugging Face. Developed a multi-agent agentic system for database engineering. Owned LLMops/MLOps, DevOps, and observability end-to-end.
- 2023-2024 Gradient.AI Machine Learning Engineer
Designed and shipped serverless fine-tuning and inference API services for large language models (LLMs), enabling elastic scaling in a multi-tenant environment.
- 2019-2022 Meta Platforms Inc Software Developer
Built machine-learning infrastructure for payments fraud detection. Maintained the HHVM open-source branch and Hack standard libraries, ported HHVM to multiple Linux distributions, and built its CI pipeline.
- 2015-2019 Thoughtworks Inc Lead / Senior Consultant
Delivered product line engineering consulting for Huawei and led architecture for Scala projects for RCA. Founded and maintained open-source projects including Binding.scala and DeepLearning.scala.
- 2006-2014 Various Game Companies Game Developer / Tech Lead
Developed and led teams to build multiple game engines, mobile/web/social games, and 3D applications. Key contributions include engine architecture, development tools, CI/CD pipelines, and open-source libraries.

Open-source projects

- 2021-2022 HHVM C++/Rust/OCaml
is Meta's virtual machine to run Hack. I ported the open source branch of HHVM to NixOS and created CI pipeline for building, testing, and packaging HHVM.
- since 2018 Dsl.scala Scala
is a framework to extend Scala language features in libraries. I also ported this framework to Haskell as Control.Dsl.
- since 2016 DeepLearning.scala and Compute.scala Scala
are libraries for creating complex neural networks.
- since 2015 Binding.scala Scala
is a reactive templating language for creating web and desktop GUI.
- 2015-2016 Microbuilder Haxe / Scala / Java / JavaScript
is a toolkit to build system across micro-services implemented in various languages communicating via RESTful JSON API.
- since 2012 haxe-continuation Haxe
provides async/await syntax for Haxe.
- since 2010 protoc-gen-as3 ActionScript3
is a Protocol Buffers plugin for ActionScript 3.
- since 2007 other 100+ of open-source projects

Publications

- 2025 CE-U: Cross Entropy Unlearning first author
- 2017 Mondic Deep Learning first author
- 2017 Ad-hoc polymorphic delimited continuations first author
- 2009 Adobe Flex 3: Training from the Source translator

Talks

- 2017 Deeplearning.scala - Thoughts and practices on OSS AI framework DBTC
- 2017 Monadic Deep Learning Strange Loop
- 2017 Neural Networks & Functional Programming ArchData Summit · Beijing
- 2016 More than Async QCon · Beijing
- 2015 Specific-domain extension in an universal language ECUG Con

Education

- 2002-2006 Southwest University of Political Science and Law Bachelor Of Laws